

Lecture 1.2

AI Introduction

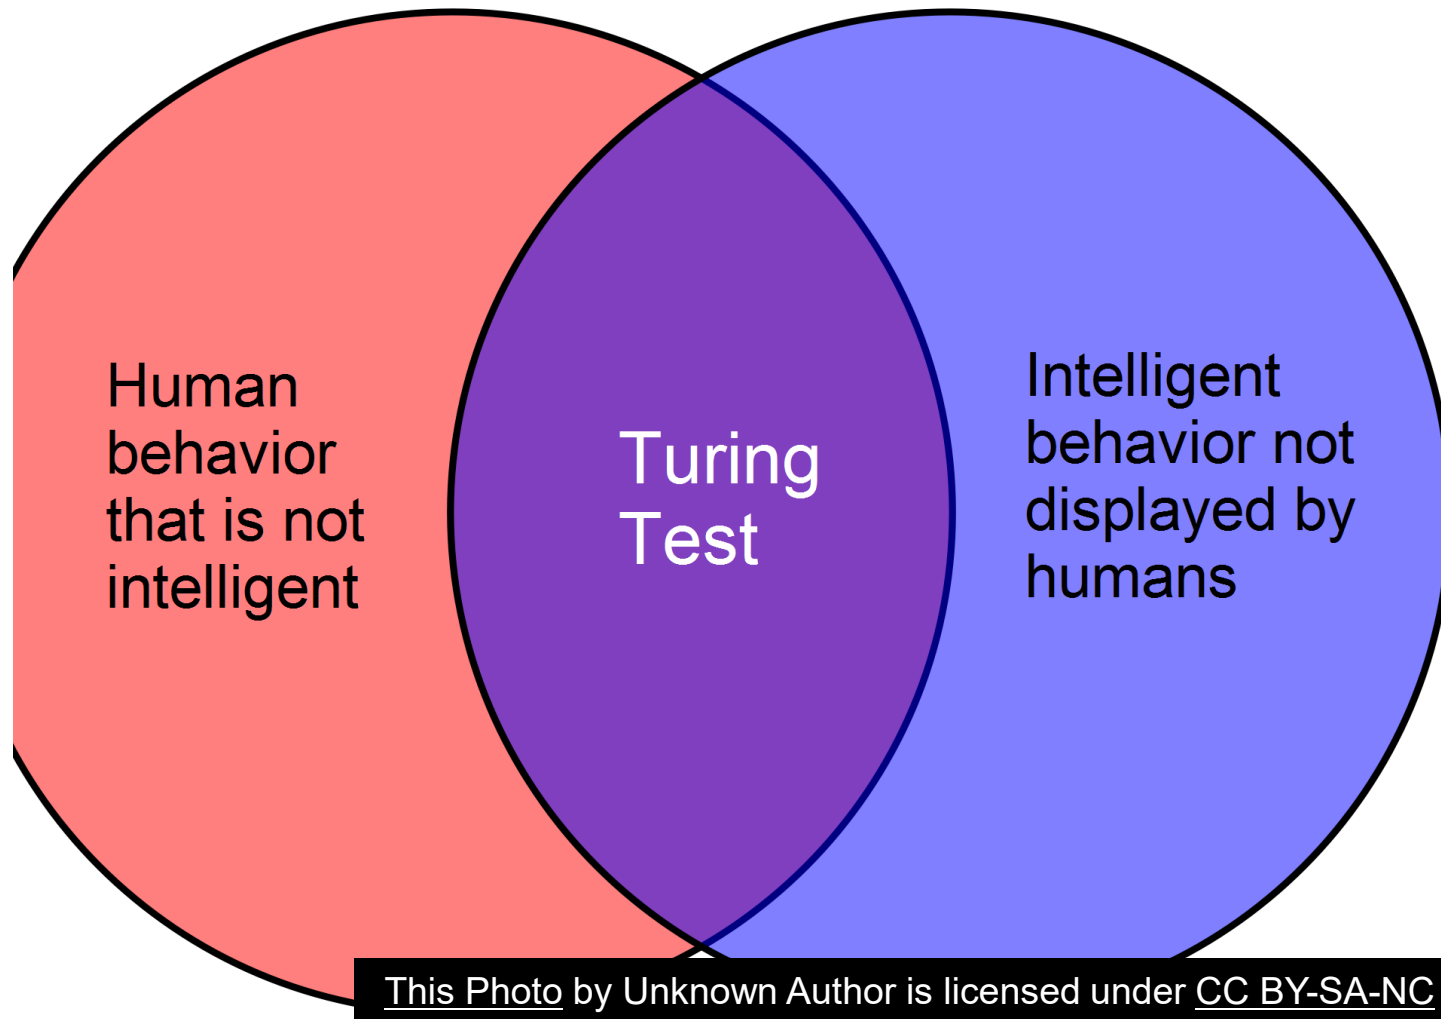
Topics

- Introduction
- What is AI
- History of AI
- NVIDIA Jetson Developer Kits
- Deep Learning

Learning Objectives

- Define artificial intelligence (AI)
- List examples of applications of AI
- Describe the different types of hardware implementations for AI computing
- Describe the difference between CPU and GPU processing
- Write and execute code and compare differences between CPU and GPUs

Human behavior Intelligent behavior



Artificial Intelligence (AI)



What is AI?

Artificial Intelligence is the emulation of human or biological behavior

Real intelligence that is generated artificially

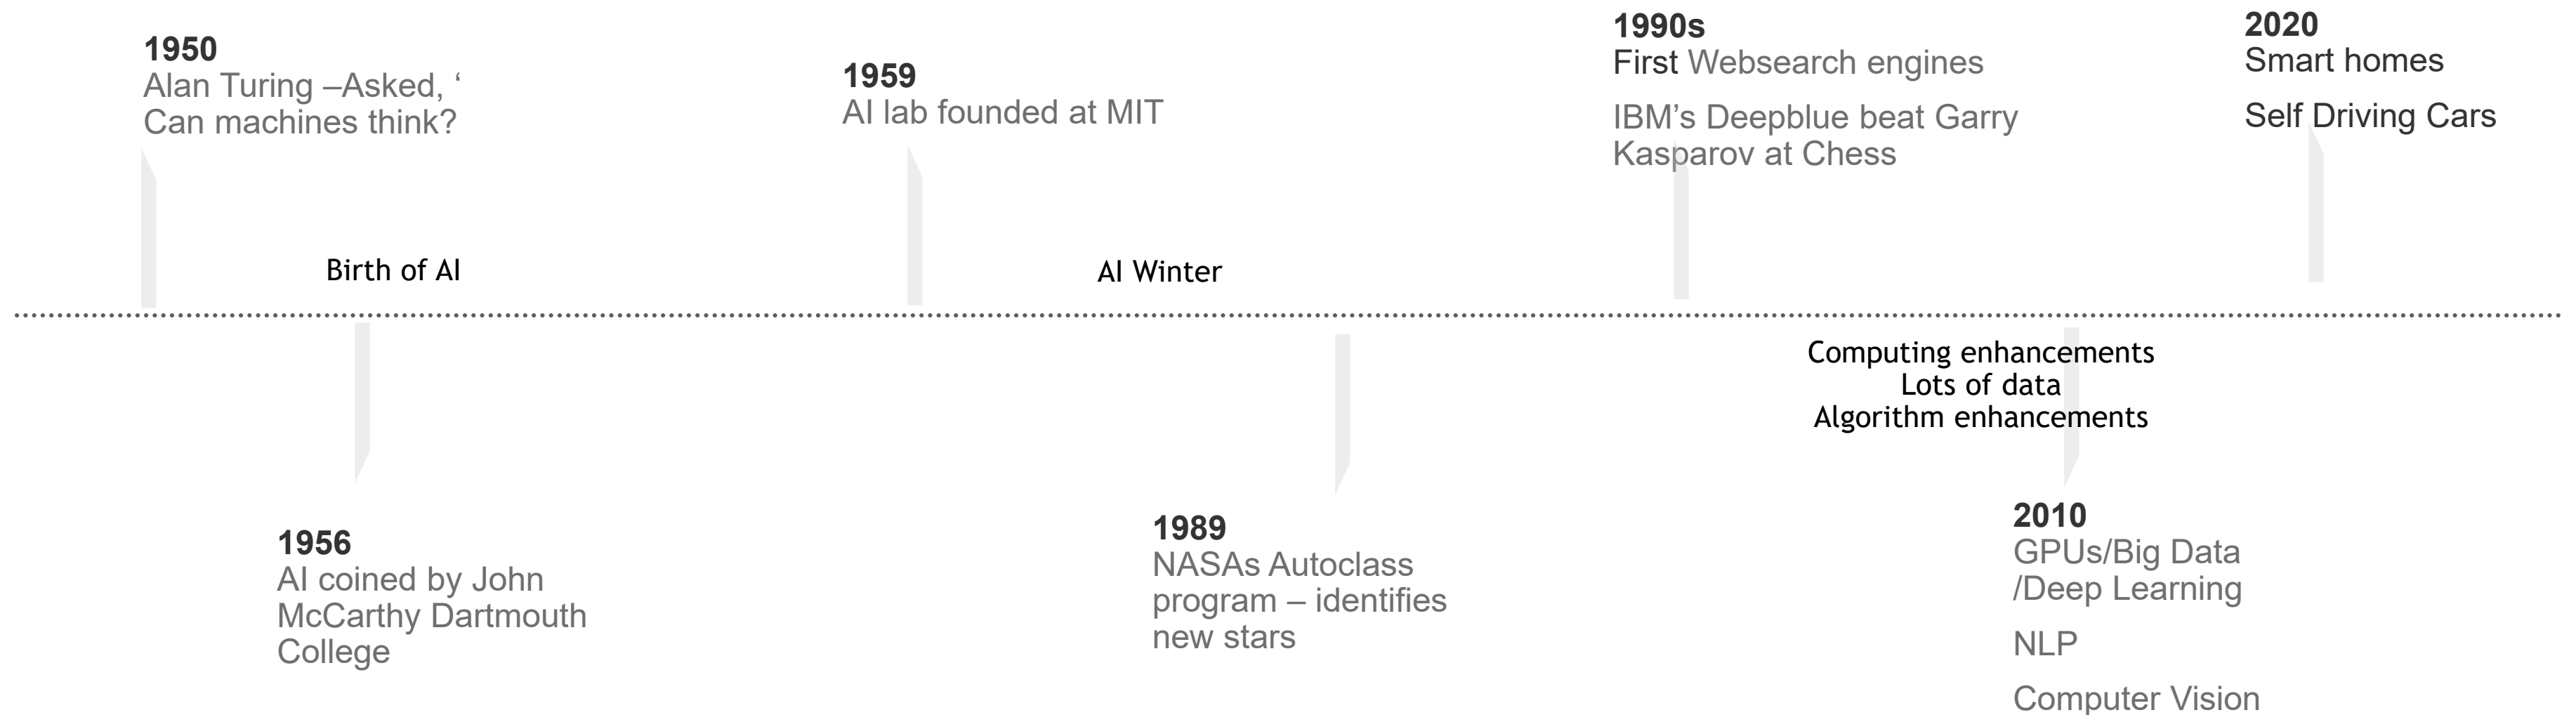
Derived from study of computational agents that act intelligently (<http://artint.info/2e/html/ArtInt2e.Ch1.S1.html>)



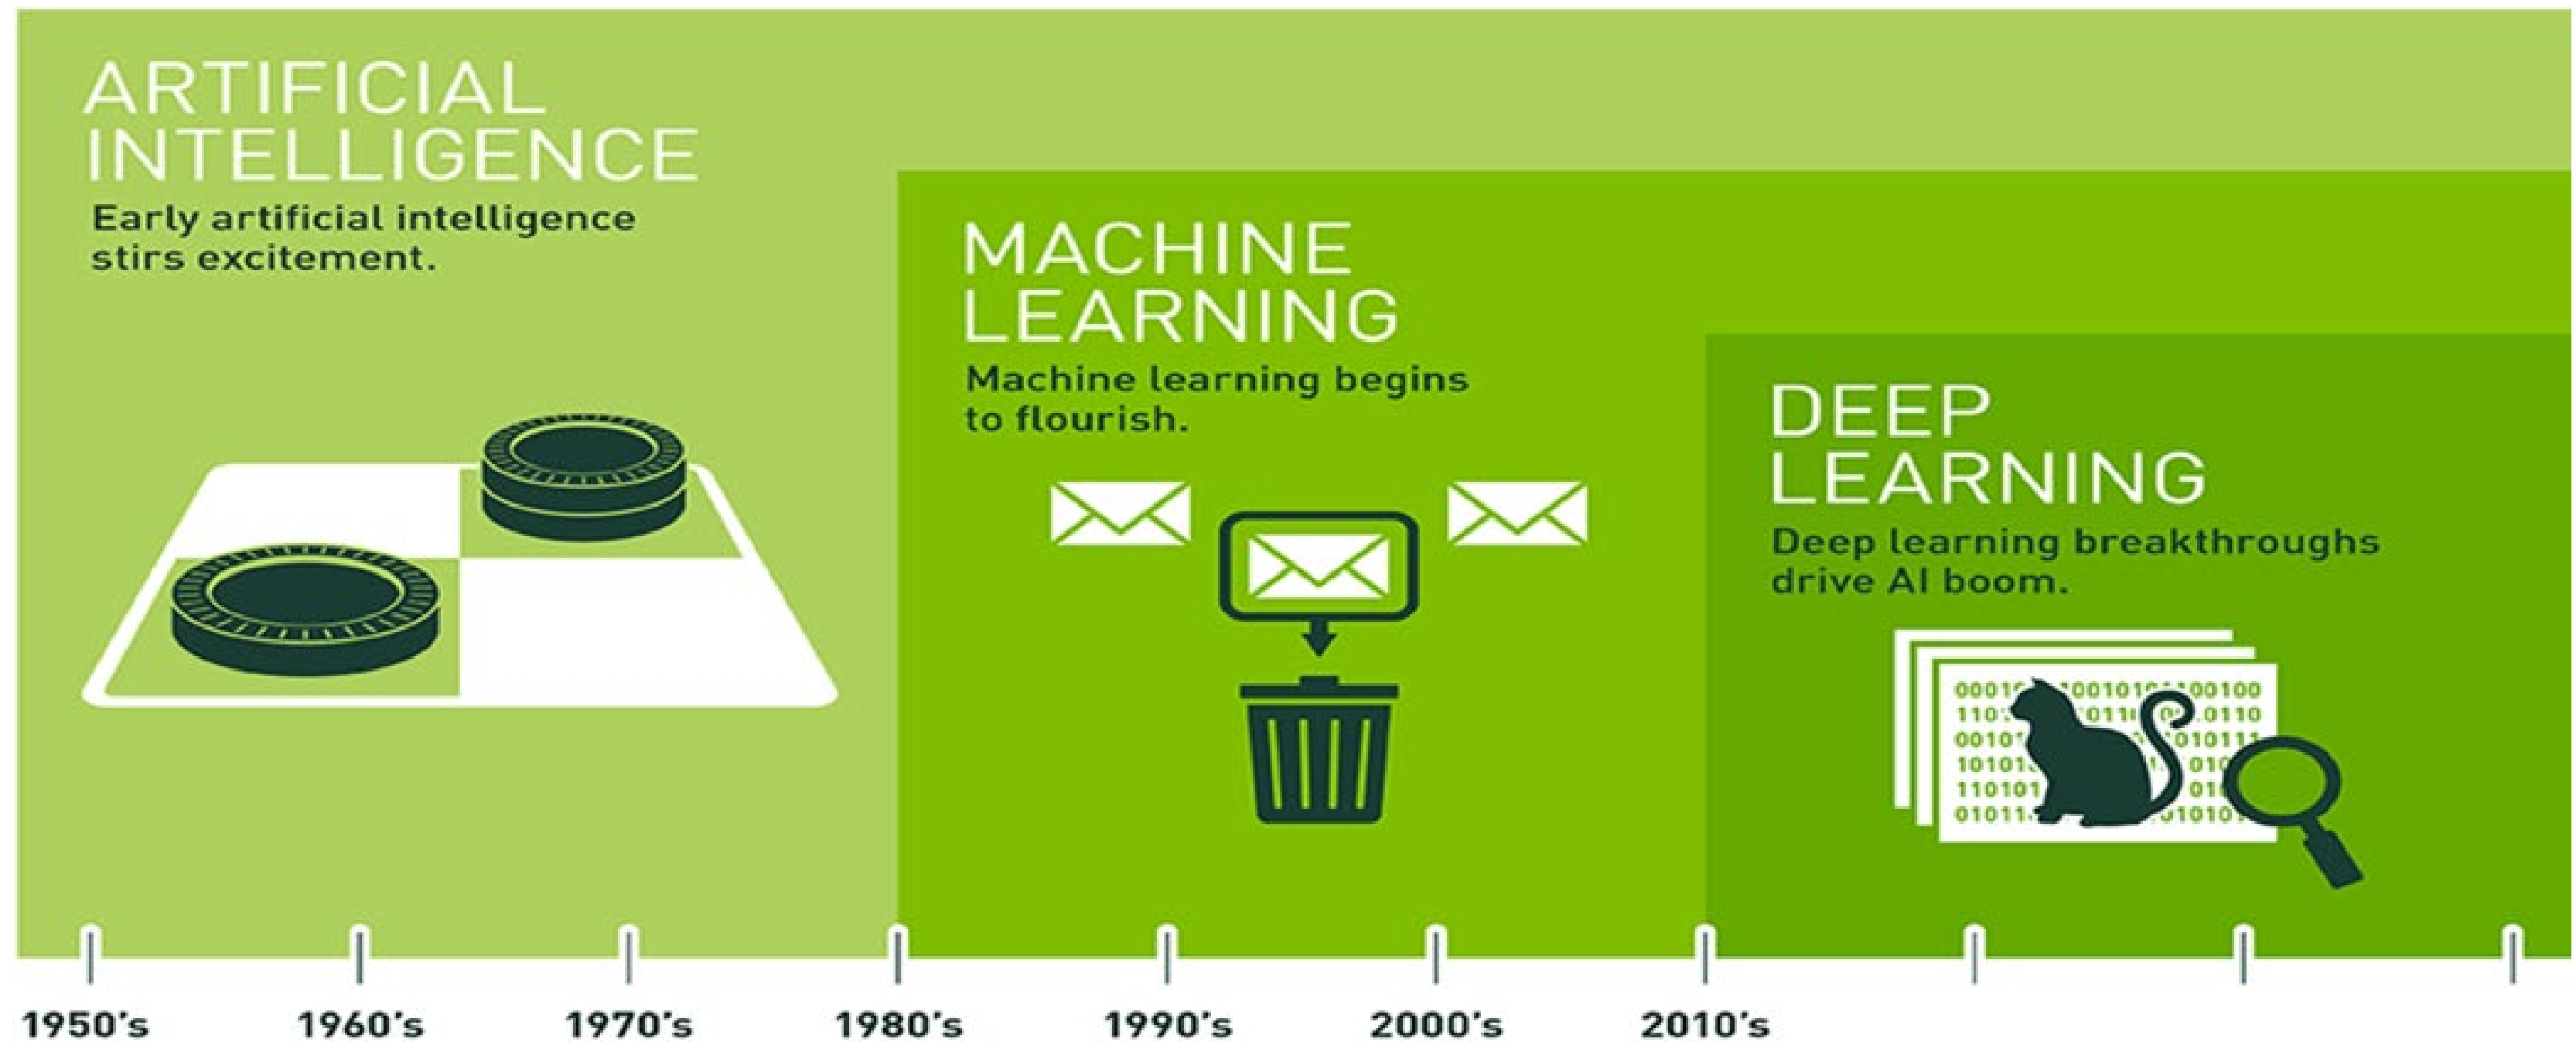
Evolution of AI


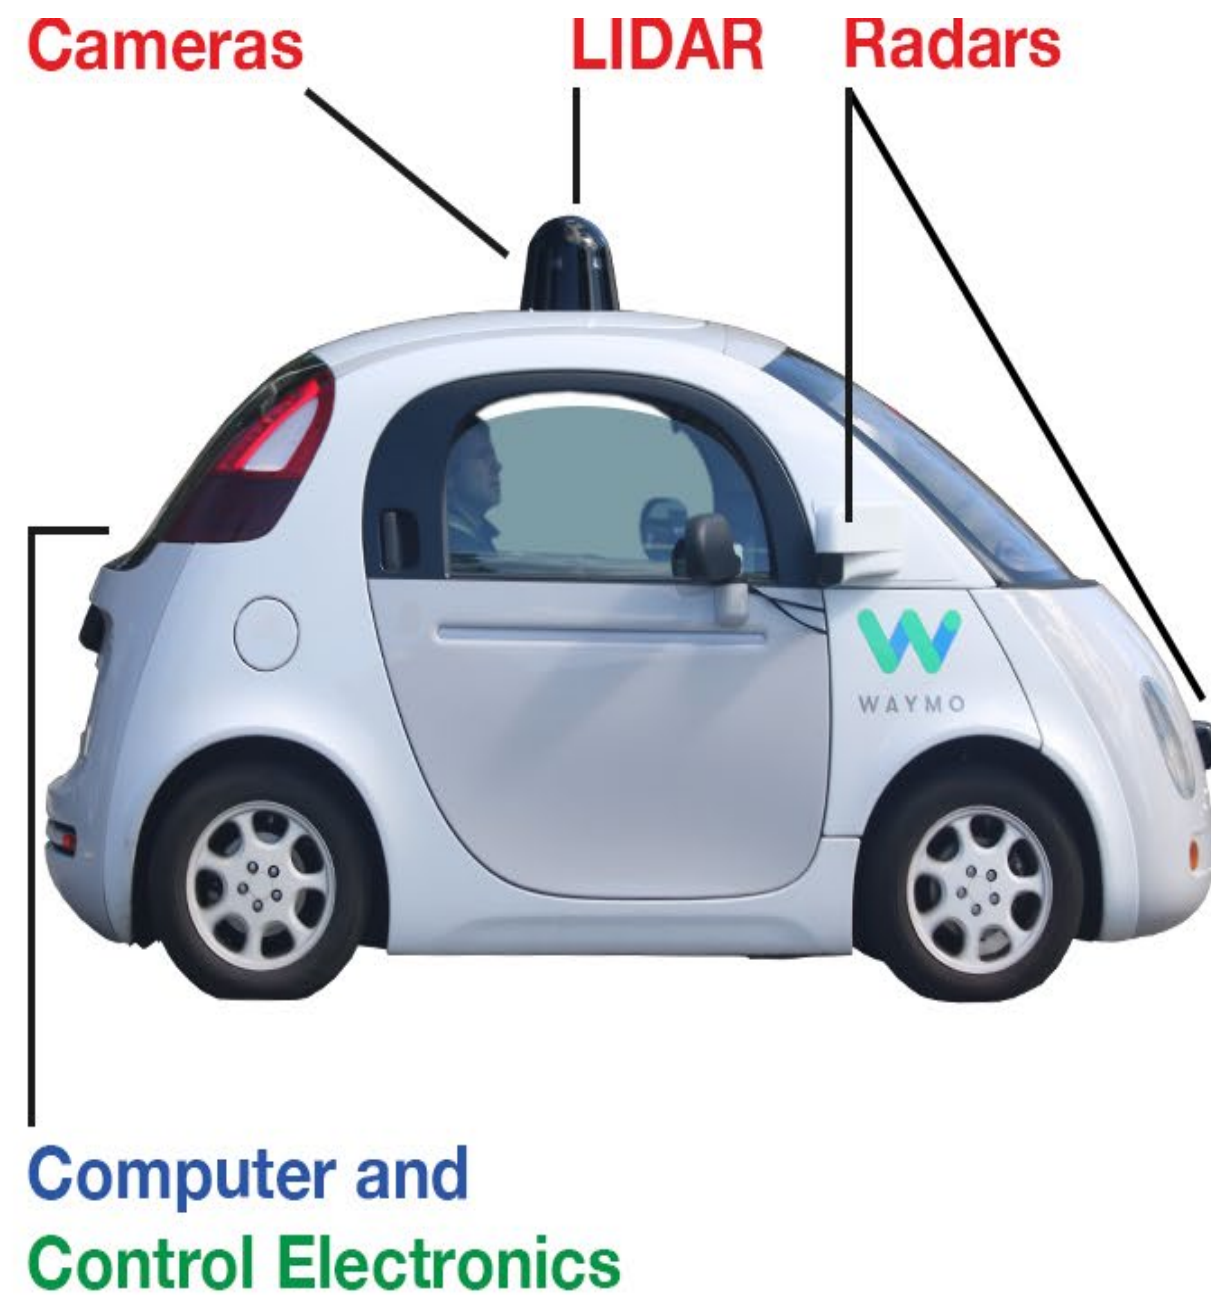
History of AI

Hobbes (1588-1679) – thinking was symbolic reasoning like talking out loud or working out an answer with pencil and paper



Evolution of AI





This Photo by Unknown Author is licensed under [CC BY-SA-NC](#)


Applied AI

Put AI into practice for real world applications

Real World Applications of AI

- Self Driving Cars
- Facial Recognition for Identification
- Deposit Checks through Mobile Apps
- Robotic Floor Cleaners
- Network Intrusion
- Loan risk models
- Recommendation engines for shopping product and movie recommendations
- Email spam detection

THE BIG BANG IN MACHINE LEARNING



DNN



GPU



BIG DATA

Framework and Hardware for Class

Much of what we will do in this course involves using Deep Learning to process images and text.

Deep Learning is computationally intensive and benefits from using GPUs to run some of those operations in parallel, especially for Image and Natural Language processing.

Rather than require expensive GPU laptops, we will use the Jetson Nano

We will use Python to code applications that use CUDA enabled libraries and deep learning framework

GPU Programming Languages

Numerical analytics ►	MATLAB, Mathematica, LabVIEW
Fortran ►	CUDA Fortran
C ►	CUDA C
C++ ►	CUDA C++
Python ►	PyCUDA, OpenCV CUDA, Numba, NumbaPro, RAPIDS
F# ►	Alea.cuBase

NVIDIA CUDA

Development environment and ecosystem to enable GPU applications

Integrates with programming languages, GPU libraries and deep learning frameworks

<https://developer.nvidia.com/cuda-toolkit>

Deep Learning Concepts

Deep Learning – Machine Learning that leverages Neural Networks with many layers

Reinforcement Learning – reward-based algorithm with the intent that actions will result in greater cumulative reward https://en.wikipedia.org/wiki/Reinforcement_learning

Federated Learning – machine learning technique enables the training of models across devices without exchanging the data https://en.wikipedia.org/wiki/Federated_learning

Deep Learning – Neural Networks

AlexNet <https://en.wikipedia.org/wiki/AlexNet> – Deep Convolutional Neural Network

Resnet https://en.wikipedia.org/wiki/Residual_neural_network – Residual Neural Network

Numerous types of Neural Networks. Helpful info on neural network architecture can be found at <https://www.asimovinstitute.org/author/fjodorvanveen/>

Popular GPU Frameworks

CAFFE <https://caffe.berkeleyvision.org/>

MXNet <https://mxnet.apache.org/versions/1.7.0/>

Tensorflow <https://www.tensorflow.org/>

PYTORCH <https://pytorch.org/>

Python – GPU Enabled Libraries

Tensorflow GPU <https://www.tensorflow.org/install/gpu>

Numba <http://numba.pydata.org/> highly performance optimized C code that can be run from within Python that is designed to work with Numpy arrays and functions.

PyCuda <https://pypi.org/project/pycuda/>

RAPIDS <https://rapids.ai/> - new GPU data science library that is modeled to provide similar look and feel of Panda and Scikit Learn libraries. It is currently not supported on the Nano

OpenCV CUDA <https://opencv.org/platforms/cuda/>