

Workflow Application in Deep Learning: Custom Datasets with PyTorch Artificial Intelligence

Zool Hilmi Ismail
Tokyo City University

“What is a custom dataset?”

“I’ve got my own dataset, can I build a model with PyTorch to predict on it?”

Yes.

PyTorch Domain Libraries

“Is this a photo of pizza, steak or sushi?”



TorchVision

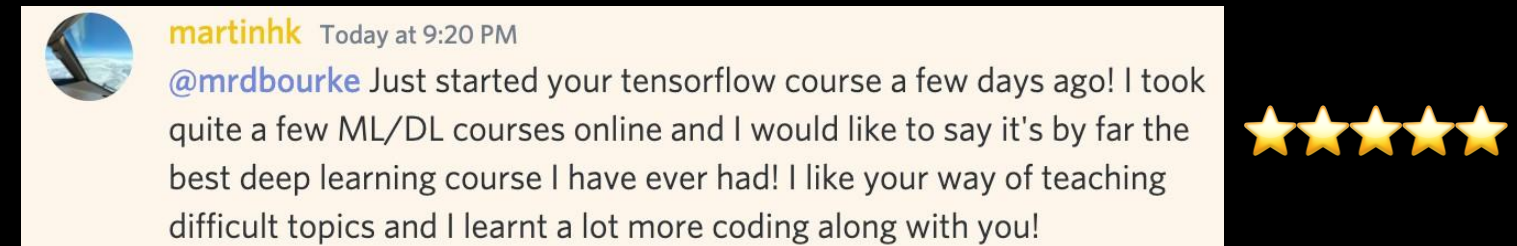
“What song is playing?”



TorchAudio

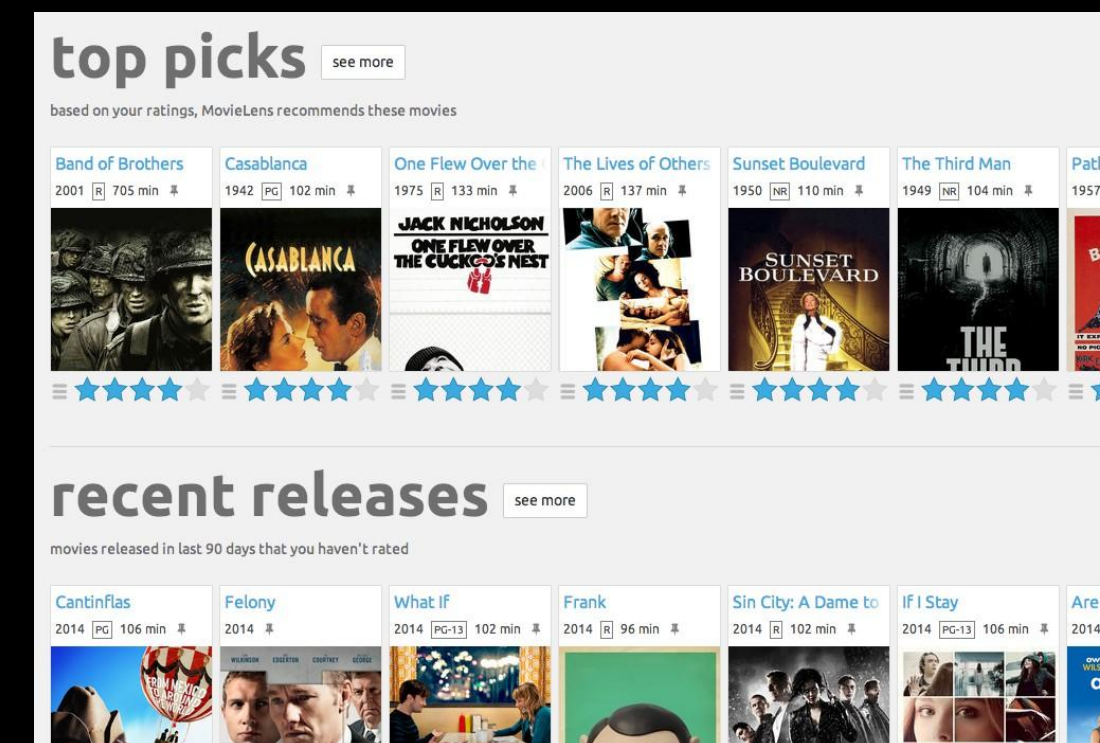
Different domain libraries contain data loading functions for different data sources

“Are these reviews positive or negative?”



TorchText

“How do we recommend similar products?”



TorchRec

Source: movielens.org


PyTorch Domain Libraries

Problem Space	Pre-built Datasets and Fuctions
Vision	<u>torchvision.datasets</u>
Text	<u>torchtext.datasets</u>
Audio	<u>torchaudio.datasets</u>
Recommendation system	<u>torchrec.datasets</u>
Bonus	<u>TorchData</u> *

*TorchData contains many diKerent helper functions for loading data and is currently in beta as of April 2022.

What we're going to build

FoodVision Mini



We're going to write code to load images of food (our own custom dataset for FoodVision Mini)

`torchvision.transforms`
`torch.utils.data.Dataset`
`torch.utils.data.DataLoader`

`torchmetrics`

`torch.save`
`torch.load`



1. Get data ready
(turn into tensors)



2. Build or pick a
pretrained model
(to suit your problem)



3. Fit the model to the
data and make a
prediction



4. Evaluate the model



5. Improve through
experimentation



6. Save and reload
your trained model



2.1 Pick a loss function & optimizer



2.2 Build a training loop

`torch.optim`

`torch.nn`
`torch.nn.Module`
`torchvision.models`

`torch.utils.tensorboard`

What we're going to cover

- Getting a **custom dataset** with PyTorch
- **Becoming one with the data** (preparing and visualising)
- **Transforming data** for use with a model
- **Loading custom data** with pre-built functions and custom functions
- Building **FoodVision Mini** to classify 🍕 🍷 🍣 images
- Comparing models with and without **data augmentation**
- **Making predictions** on custom data

(we'll be cooking up lots of code!)

How:



Let's code!



Daniel Bourke

@mrdbourke



“If I had 8 hours to build a machine learning model, I’d spend the first 6 hours preparing my dataset.”

- Abraham Lossfunction

12:35 PM · Nov 4, 2021 · Twitter Web App

Standard image classification data format

**Your own data format
will depend on what
you're working**

```
pizza_steak_sushi/ # <- overall dataset folder
train/ # <- training images
  pizza/ # <- class name as folder name
    image01.jpeg
    image02.jpeg
    ...
  steak/
    image24.jpeg
    image25.jpeg
    ...
  sushi/
    image37.jpeg
    ...
test/ # <- testing images
  pizza/
    image101.jpeg
    image102.jpeg
    ...
  steak/
    image154.jpeg
    image155.jpeg
    ...
  sushi/
    image167.jpeg
    ...
```

**The premise remains:
write code to get your
data into tensors for
use with PyTorch**

What is data augmentation?


Looking at the same image but from different perspective(s)*. To artificially increase the diversity of a dataset.



Original



Rotate



Shift

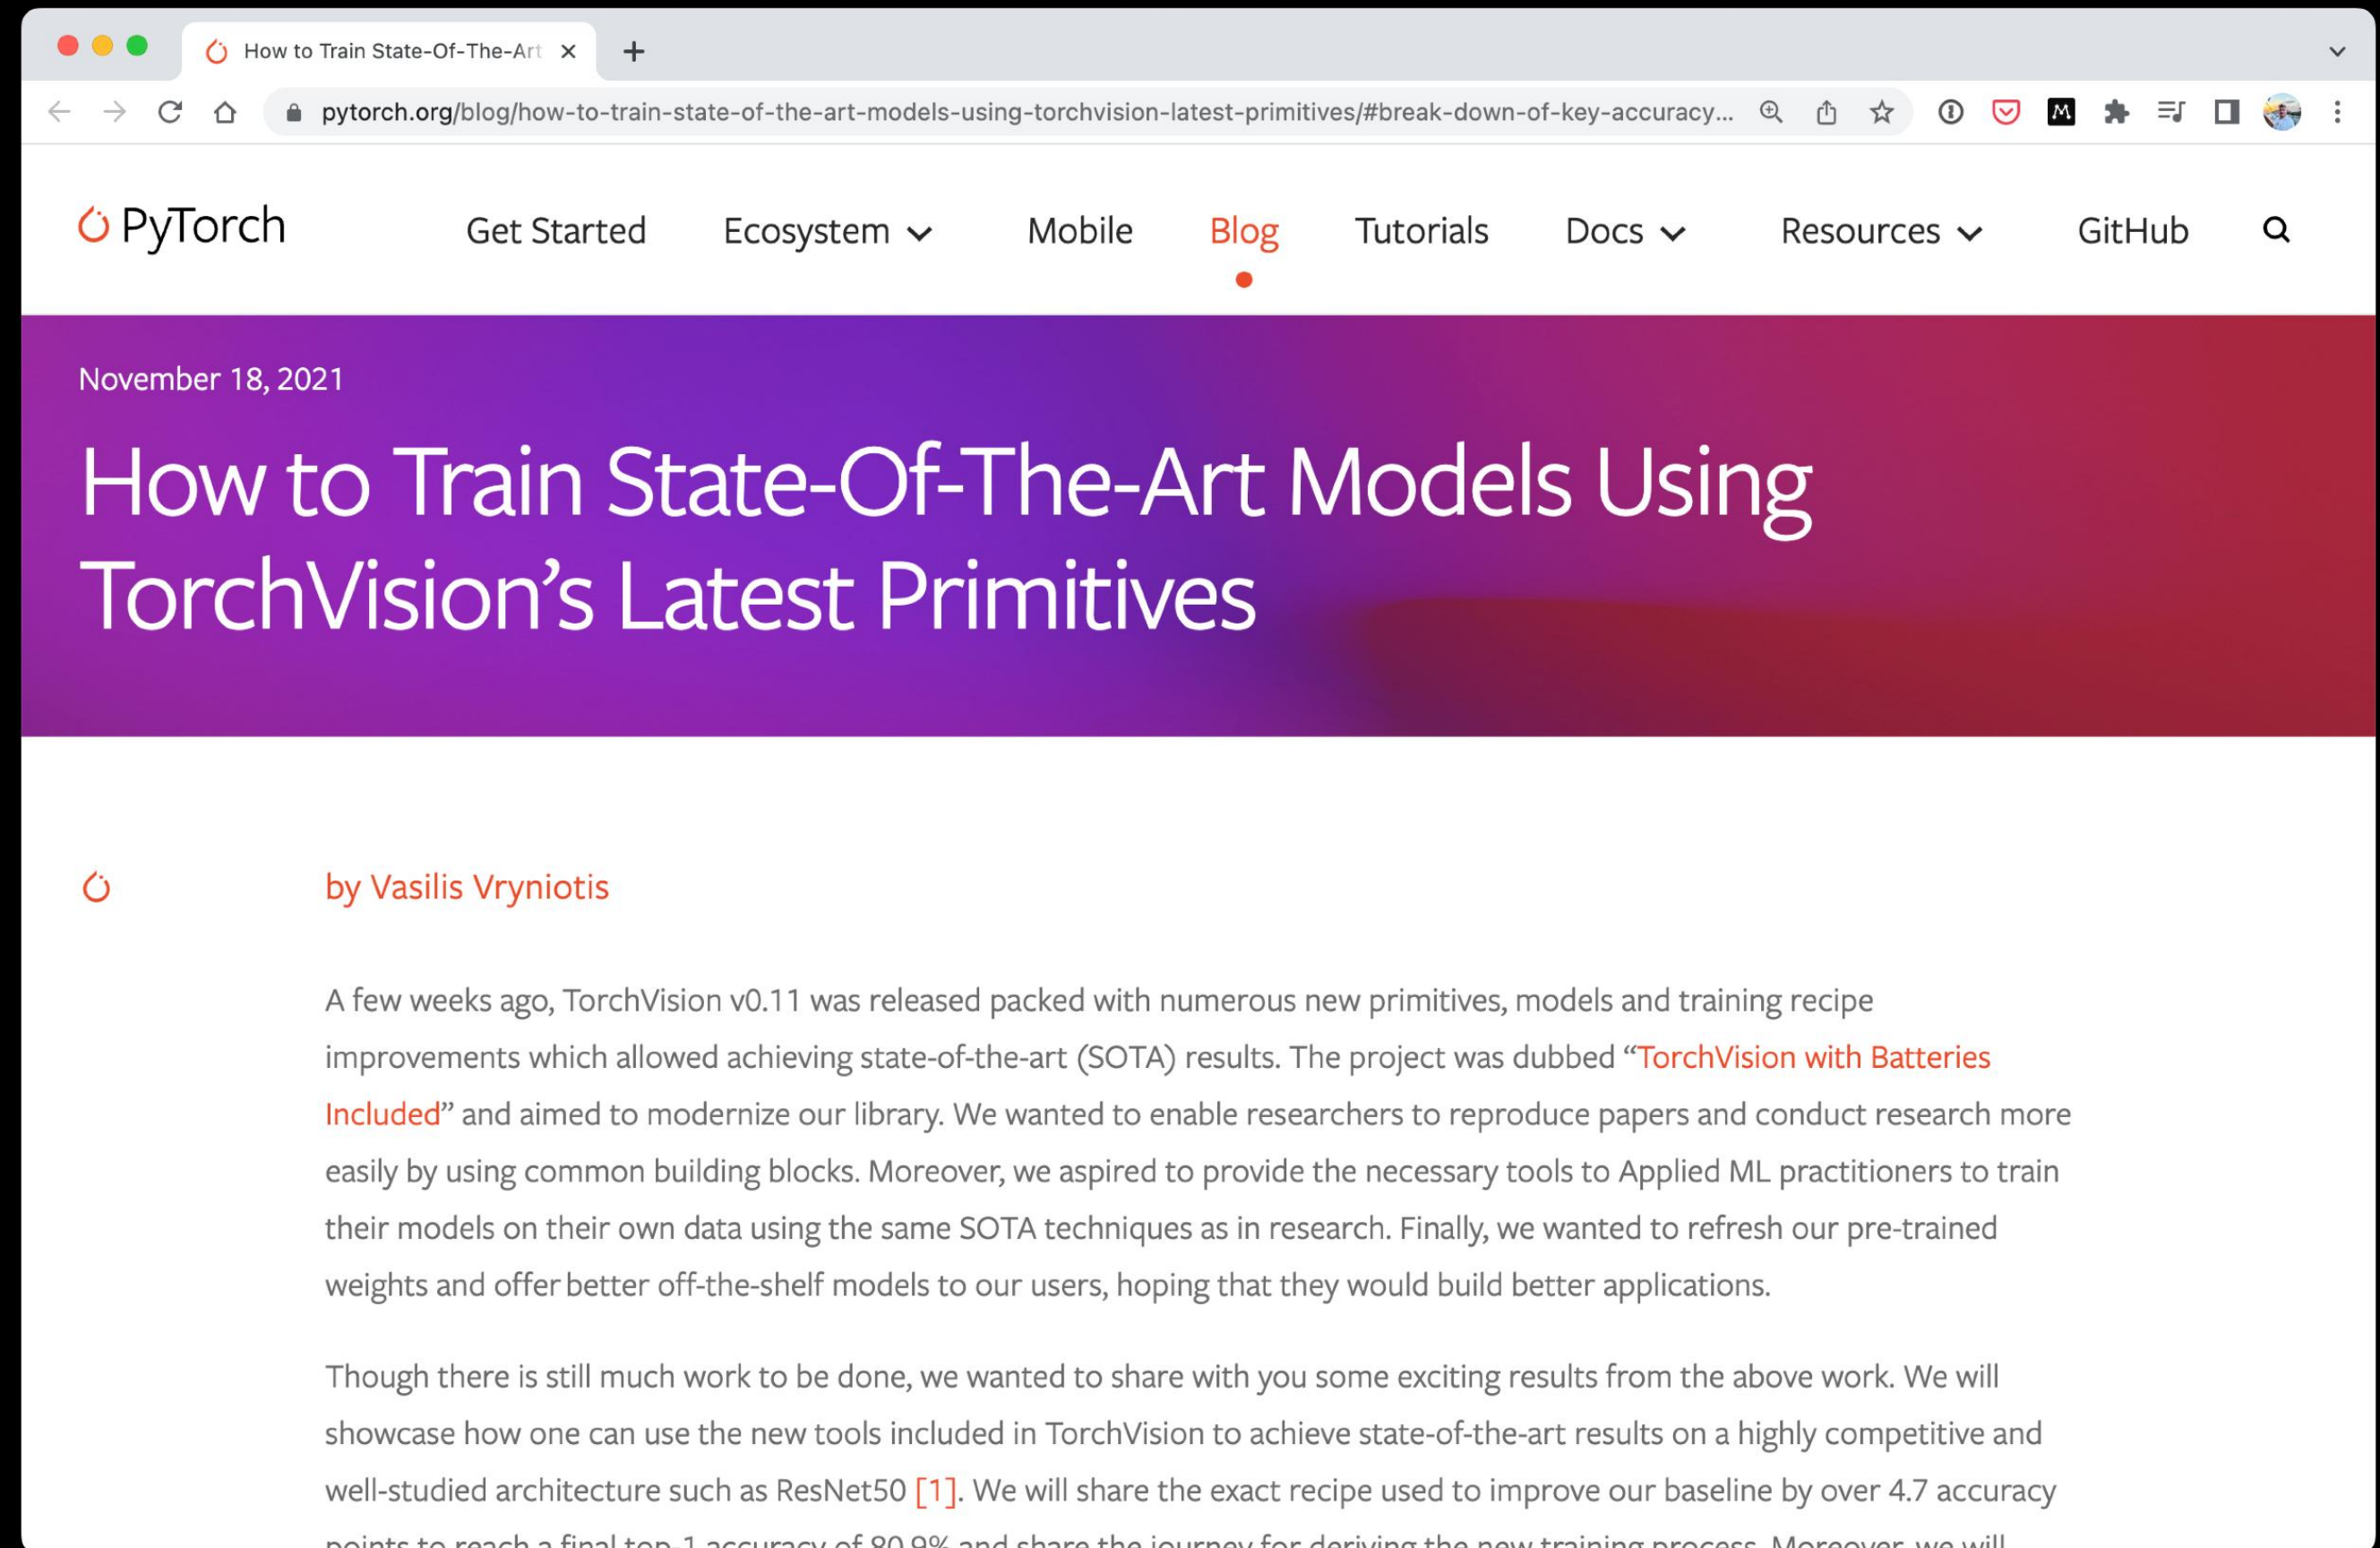


Zoom

***Note:** There are many more different kinds of data augmentation such as, cropping, replacing, shearing. This slide only demonstrates a few.

PyTorch State of the Art Recipe


Research comes out often on how best to train models, state-of-the-art (SOTA) methods are always changing).



Source: Training state-of-the-art computer vision models with torchvision from the PyTorch blog.

Loss curves

(a way to evaluate your model's performance over time)



Dealing with overfitting

Method to improve a model (reduce overfitting)	What does it do?
Get more data	Gives a model more of a chance to learn patterns between samples (e.g. if a model is performing poorly on images of pizza, show it more images of pizza).
Data augmentation	Increase the diversity of your training dataset without collecting more data (e.g. take your photos of pizza and randomly rotate them 30°). Increased diversity forces a model to learn more generalisation patterns.
Better data	Not all data samples are created equally. Removing poor samples from or adding better samples to your dataset can improve your model's performance.
Use transfer learning	Take a model's pre-learned patterns from one problem and tweak them to suit your own problem. For example, take a model trained on pictures of cars to recognise pictures of trucks.
Simplify your model	If the current model is already overfitting the training data, it may be too complicated of a model. This means it's learning the patterns of the data too well and isn't able to generalize well to unseen data. One way to simplify a model is to reduce the number of layers it uses or to reduce the number of hidden units in each layer.
Use learning rate decay	The idea here is to slowly decrease the learning rate as a model trains. This is akin to reaching for a coin at the back of a couch. The closer you get, the smaller your steps. The same with the learning rate, the closer you get to <u>convergence</u> , the smaller you'll want your weight updates to be.
Use early stopping	<u>Early stopping</u> stops model training *before* it begins to overfit. As in, say the model's loss has stopped decreasing for the past 10 epochs (this number is arbitrary), you may want to stop the model training here and go with the model weights that had the lowest loss (10 epochs prior).

Dealing with underfitting

Method to improve a model (reduce underfitting)	What does it do?
Add more layers/units to your model	If your model is underfitting, it may not have enough capability to <i>learn</i> the required patterns/weights/representations of the data to be predictive. One way to add more predictive power to your model is to increase the number of hidden layers/units within those layers.
Tweak the learning rate	Perhaps your model's learning rate is too high to begin with. And it's trying to update its weights each epoch too much, in turn not learning anything. In this case, you might lower the learning rate and see what happens.
Train for longer	Sometimes a model just needs more time to learn representations of data. If you find in your smaller experiments your model isn't learning anything, perhaps leaving it train for a more epochs may result in better performance.
Use transfer learning	Take a model's pre-learned patterns from one problem and tweak them to suit your own problem. For example, take a model trained on pictures of cars to recognise pictures of trucks.
Use less regularization	Perhaps your model is underfitting because you're trying to prevent overfitting too much. Holding back on regularization techniques can help your model fit the data better.

Predicting on custom data

(3 things to make sure of...)

Custom image



`[[0.31, 0.62, 0.44...],
[0.92, 0.03, 0.27...],
[0.25, 0.78, 0.07...],
...]`

`torch.float32`

1. Data in right datatype

Is the model on the GPU?



2. Data on same device as model

Original

`Shape = [64, 64, 3]`

Add batch dimension & rearrange if needed

`Shape = [None, 64, 64, 3] (NHWC)`

`Shape = [None, 3, 64, 64] (NCHW)`

Same as model input

3. Data in correct shape